

TABLE OF CONTENTS

Abstract	2
1. Introduction	2
2. Serverless Computing Architecture and Challenges	3
2.1 Basic Concepts of Serverless Computing	3
2.2 Importance and Challenges of Function Composition	4
2.3 Performance Bottlenecks and Existing Optimization Strategies	5
3. Related Work	5
3.1 Static Scheduling Methods	6
3.2 Dynamic Scheduling Methods	6
3.3 Limitations of Existing Optimization Algorithms	7
4. Design and Analysis of Function Composition Algorithms	8
4.1 Algorithm Design	8
4.2 Performance Analysis	9
4.3 Complexity Analysis of the Algorithm	10
5. Experiments and Evaluation	10
5.1 Experimental Setup	10
5.2 Experimental Results and Analysis	11
5.3 Comparative Analysis and Discussion	13
6. Conclusion and Future Work	14
Conclusion	15
Acknowledgement	15
References	16

Research on Serverless Computing Based on Function Composition Algorithms

Xiaoqi Yang

Computer Science and Technology, China University of Mining and Technology
Xuzhou, China

Abstract

Serverless computing has become a critical research area in distributed and cloud computing, particularly in the optimization of function composition algorithms. With the rapid development of cloud architectures, the event-driven serverless computing model demonstrates significant potential in achieving resource elasticity and automatic scaling. This paper explores the application of function composition algorithms in serverless computing, analyzes existing composition methods, identifies performance bottlenecks, and proposes a novel algorithmic optimization strategy. Through experimental evaluation, we demonstrate the advantages of the proposed algorithm in improving computational efficiency, reducing latency, and optimizing resource allocation. The results show that the optimized function composition algorithm effectively enhances the overall performance of serverless computing platforms, especially when handling large-scale concurrent requests.

Keywords: Serverless Computing, Function Composition, Performance Optimization, Event-Driven Architecture, Scheduling Algorithms

1. Introduction

Serverless computing, as a cutting-edge cloud computing model (as shown in Figure 1), is transforming the way software applications are built and managed in this paper. It not only represents a technological innovation but also poses a profound challenge to traditional IT resource management concepts. Under this model, service providers are responsible for configuring, maintaining, and scaling servers, while developers can focus entirely on implementing business logic without worrying about underlying infrastructure issues. This shift significantly lowers the barrier to starting new projects, accelerates time-to-market for products, and through a pay-as-you-go billing model, makes cost control more flexible.

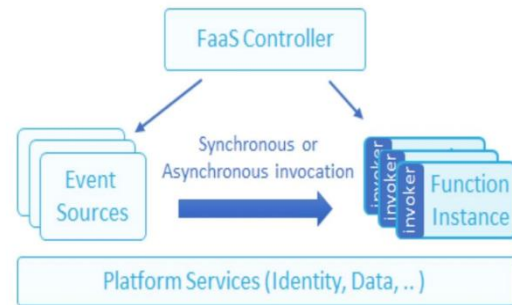


fig.1. Serverless Computing Model

In serverless architecture, Function as a Service (FaaS) is one of the core components. Each function is an independent, lightweight service unit that can be automatically triggered to execute based on actual requests. This highly decoupled design not only simplifies code structure but also supports almost infinite

horizontal scalability. However, as the complexity of enterprise applications increases, a single function often fails to meet the requirements of complex business scenarios. Therefore, how to efficiently combine multiple functions to form organized service chains or workflows while ensuring overall performance remains unaffected has become a critical issue that needs to be addressed.

Moreover, in response to the growing demand for data processing, optimizing the invocation mechanism between functions, reducing cold start times, and improving concurrent processing efficiency have also become major challenges faced by serverless architecture. To tackle these challenges, researchers and engineers are actively exploring new solutions, such as introducing asynchronous communication models, adopting more advanced caching strategies, and even utilizing containerization technology to replace traditional function execution environments under certain circumstances.

Through these technological innovations, serverless computing not only enhances the flexibility and scalability of applications but also significantly reduces operational costs for enterprises. Both startups and large enterprises can benefit from this, enjoying more efficient and economical cloud service experiences.

This study aims to analyze and optimize function composition algorithms, proposing a new efficient combination strategy. Theoretical analysis and experimental verification will be conducted to demonstrate its application value in serverless computing.

2. Serverless Computing Architecture and Challenges

Serverless computing is an emerging computational model designed to liberate developers from traditional server management and resource scheduling. Nowadays, the

application of serverless computing has become very widespread, as shown in Figure 2. The core concept of this architecture is to trigger function execution through event-driven mechanisms without the need to pre-configure or manage underlying server resources. Compared to traditional virtual machines and containerized architectures, serverless computing offers higher levels of automation, scalability, and flexibility, significantly reducing the burden of operations and maintenance. In a serverless architecture, developers only need to write business logic functions, while the platform automatically handles resource allocation, load balancing, and scaling issues. Typically, service providers charge on a pay-per-use basis according to the actual amount of computation used, optimizing resource utilization.




fig.2. Main Application Scenarios of Serverless Computing

2.1 Basic Concepts of Serverless Computing

The fundamental characteristic of the serverless computing model is "event-driven," as illustrated in Figure 3. Under this model, functions (commonly referred to as "serverless functions" or "Serverless Functions") are automatically executed in response to events, which can include HTTP requests, file uploads, database updates, message queue events, and more. This mode typically relies on the infrastructure provided by cloud computing

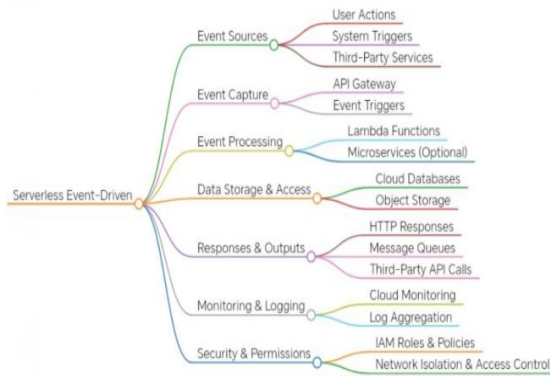


fig.3. Event-driven Model

platforms (such as AWS Lambda, Google Cloud Functions, Azure Functions, etc.) to manage the execution of functions. Unlike traditional computing resource models, in serverless computing, developers do not need to worry about the configuration, maintenance, or scaling of underlying servers; instead, they focus on writing business logic.

The advantage of this architecture lies in its "serverless" nature, where users do not pay for reserved computing resources. Users only pay for the actual execution time of their functions and the resources consumed, thereby reducing the cost of infrastructure. Serverless architecture offers extremely high scalability, automatically adjusting the amount of computing resources based on changes in request volume. It can automatically increase resources during peak times and reduce them during off-peak periods. As a result, the system can efficiently handle dynamic load variations, avoiding the risk of over-provisioning resources to cope with sudden traffic spikes in traditional architectures.

2.2 Importance and Challenges of Function Composition

In serverless computing, function composition is a common practice, especially in complex business processes where the output of one function often serves as the input

for another. For example, when a user uploads a file, it can trigger one function to process the file, followed by another function to save the results to a database. This combination and orchestration of functions are a core component of serverless computing.

The implementation of function composition requires efficient event passing mechanisms and coordination between functions. As shown in Figure 4, in some complex application scenarios, multiple functions may need to work together, such as in image processing, video encoding and decoding tasks, etc. The dependencies and data transfer methods between these functions need to be carefully designed. A common challenge is ensuring the efficient scheduling of these functions to avoid performance bottlenecks or resource wastage caused by improper execution order. Particularly in high-concurrency situations, ensuring the correct sequence and accuracy of function calls is a significant concern.



fig.4. Function Scheduling under High Concurrency

Moreover, function composition involves more than just code-level coordination; it also encompasses performance optimization and resource management. For instance, the combination of multiple functions can introduce unnecessary latency, especially when there is frequent network communication and complex function call dependencies, leading to a noticeable decrease in system response speed. To overcome these issues, developers need to adopt asynchronous mechanisms, event-driven architectures, and integrate techniques such as caching and message queues to reduce the latency of function calls.

2.3 Performance Bottlenecks and Existing Optimization Strategies

Despite the great convenience and flexibility offered by serverless computing, there are several performance bottlenecks. Below are some of the main performance bottlenecks and current optimization strategies:

Cold Start Latency: In serverless computing, a "cold start" occurs when a function has not been invoked for a long time, and the cloud platform needs to initialize the necessary resources and load the code, resulting in latency. This cold start delay can affect system response time, particularly in applications requiring low-latency responses (As shown in Figure 5). To optimize cold start issues, developers often use a "keep warm" strategy, triggering functions periodically to prevent prolonged inactivity. Another approach is to pre-deploy frequently used functions in memory using containerization technology to reduce initialization time.

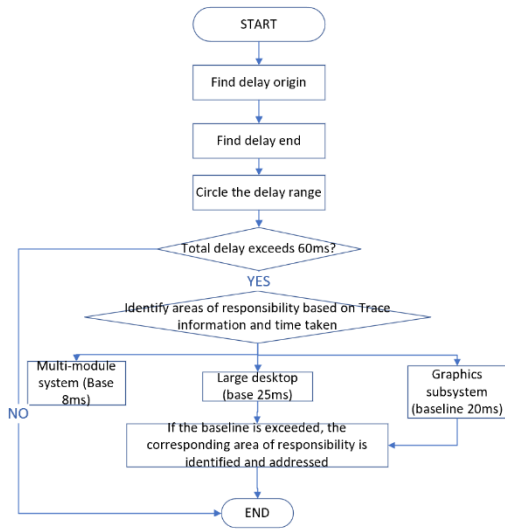


fig.5. Cold Start Response Delay Problem Processing Flow

Function Execution Time Limits: Many cloud platforms impose limits on the execution time of individual serverless functions (for example, the maximum execution time for AWS Lambda is 15 minutes). When a function's processing takes longer, the task may

need to be split, or an asynchronous execution model can be adopted to ensure timely completion. Additionally, platforms offer options to increase parallelism, handling large volumes of requests in parallel to reduce overall response time.

Resource Management and Scheduling Issues: Resource scheduling and management in serverless computing can become complex under high loads. In some cases, uneven computational demands or large data processing volumes can lead to over-allocation or under-allocation of resources. To address this, developers can implement load balancing and auto-scaling strategies, dynamically adjusting resource allocation through monitoring and log analysis. Furthermore, leveraging container technology ensures efficient resource utilization and reduces idle resource situations.

Overhead of Cross-Service Invocations: In serverless computing, cross-service communications (such as function-to-function calls and external API calls) can introduce additional latency. When designing function compositions, it is advisable to minimize the number of inter-function calls and dependencies, using batch operations or distributed caching to reduce frequent data exchanges.

Overall, while serverless computing faces certain challenges, these can be effectively mitigated through proper architectural design, optimization strategies, and technical means. As cloud computing technology continues to evolve, the architecture and optimization strategies for serverless computing will also improve, providing developers with more efficient solutions.

3. Related Work

The problem of function composition has been studied for a long time in the fields of

distributed computing and cloud computing. With the rise of serverless computing architectures, increasing research has focused on how to efficiently schedule and manage the execution of multiple functions. Function composition involves organizing and scheduling different functions to maximize resource utilization, minimize latency, and enhance overall system performance. Current research primarily focuses on task scheduling, optimizing the execution order of functions, and resource allocation strategies, especially in distributed computing environments. As shown in Figure 6, these studies include both static and dynamic scheduling methods and propose various optimization strategies to address the challenges of function composition and scheduling.

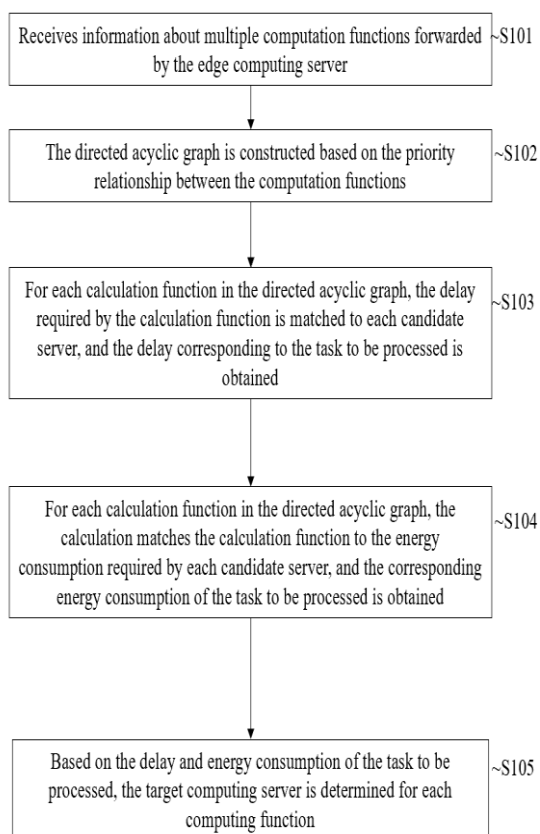


fig.6. Task Scheduling Methods in Serverless Computing

3.1 Static Scheduling Methods

Static scheduling methods rely on predefined task execution information, such as dependencies between tasks, resource requirements, and estimated execution times. Using this static information, scheduling algorithms can plan ahead and determine the execution order and resource allocation for each function before the tasks begin, thereby optimizing the overall system performance. A common goal of static scheduling is to minimize the total execution time (i.e., makespan) or maximize resource utilization.

Static scheduling methods typically use task dependency graphs (such as DAGs, Directed Acyclic Graphs) to represent dependencies between different functions. In this model, each function is represented as a node, and dependencies between functions are represented as edges. Based on this, scheduling algorithms can use classical scheduling strategies, such as greedy algorithms, shortest job first, longest job first, etc., to assign tasks. The advantage of static scheduling is its deterministic nature, allowing the system to know exactly how to proceed before execution.

However, static scheduling has limitations, especially in highly dynamic cloud environments. Since the resource state, load conditions, and function execution times can change at different points in time, static scheduling methods often fail to adapt to these dynamic changes, leading to issues such as uneven resource allocation and load imbalance. As a result, static scheduling methods may not fully utilize resources, reducing system efficiency.

3.2 Dynamic Scheduling Methods

Dynamic scheduling methods are more flexible and do not rely on static task information. Instead, they adjust based on the real-time state of tasks. Dynamic scheduling typically combines load monitoring, system

performance analysis, and resource scheduling algorithms to dynamically adjust the execution order and resource allocation of tasks during function execution. A key feature of dynamic scheduling is its ability to adapt scheduling strategies based on current system load, network conditions, and resource availability, thus addressing runtime changes.

A common strategy in dynamic scheduling is priority-based scheduling. This method assigns a dynamic priority to each function based on factors such as urgency, resource requirements, and other considerations, and adjusts these priorities during execution. For example, when system load is high, dynamic scheduling algorithms can defer the execution of low-priority functions and prioritize critical tasks with more computational resources. Another common strategy is fault-tolerant scheduling, which monitors the status of functions during execution and immediately takes alternative actions or reschedules if faults or performance degradation occur.

The advantage of dynamic scheduling is its high adaptability, effectively handling resource fluctuations and task delays. However, this method typically incurs higher computational overhead due to the need for real-time monitoring and decision-making, increasing system complexity. Additionally, dynamic scheduling must consider multiple factors, such as resource consumption, task dependencies, and network latency, making algorithm design and implementation more complex.

Table 1

Dynamic Scheduling	Static Scheduling
Determined at runtime	Determined at compile time
Flexible resource allocation	Fixed resource allocation
Suitable for varying workloads	Suitable for stable workloads

3.3 Limitations of Existing Optimization Algorithms

While static and dynamic scheduling methods have improved function execution and resource allocation in serverless computing to some extent, existing optimization algorithms still have several limitations.

Limitations of Static Scheduling: Static scheduling methods assume that system resources and task behaviors are known and stable. However, in actual serverless environments, resource volatility and the dynamic nature of tasks make it difficult for static scheduling to perform optimally. For example, static scheduling cannot effectively handle changes in server load and variations in task execution times, potentially leading to premature or delayed task execution, which affects system performance. Static scheduling methods have poor adaptability to dynamic loads and cannot handle time-sensitive requirements or sudden traffic surges.

Complexity of Dynamic Scheduling: While dynamic scheduling methods can adjust task scheduling based on real-time data, their algorithms are often very complex and require continuous collection and analysis of system states, which can result in significant computational overhead. Additionally, dynamic scheduling has high real-time performance requirements, and any delay in system monitoring and feedback mechanisms can miss optimal scheduling opportunities, leading to performance degradation. Dynamic scheduling also needs to manage dependencies between functions and resource contention during parallel execution, further increasing the complexity of scheduling algorithms.

Imbalanced Resource Allocation: Existing scheduling methods often assume equal resource distribution, but in practical applications, differences in resource requirements, network bandwidth limitations,

and varying quality of service (QoS) requirements make simple resource allocation strategies inadequate. Resource-intensive tasks may consume a disproportionate amount of computational resources, causing delays in other tasks and wasting system resources. Therefore, achieving balanced resource allocation while maintaining performance remains a significant research challenge.

Task Scheduling and Fault Tolerance: Many optimization algorithms do not adequately consider the risk of task failure and fault tolerance mechanisms. In practical applications, functions may encounter network failures, hardware failures, or other anomalies due to the instability of cloud platform resources. How to quickly and effectively reschedule and resume execution in such cases is a critical issue. Existing algorithms still have room for improvement in fault tolerance handling.

Overall, while existing optimization algorithms have enhanced task scheduling efficiency in serverless computing to some extent, several challenges remain unresolved. Future research needs to explore how to combine static and dynamic scheduling strategies to improve overall system performance and reduce resource waste. As serverless computing architectures continue to evolve, new algorithms and optimization methods will be proposed to address the limitations of existing approaches.

4. Design and Analysis of Function Composition Algorithms

The core objective of function composition algorithms is to optimize the overall system performance, including execution time, resource utilization, and latency, by effectively arranging and scheduling the execution of multiple functions. This section will detail a function composition algorithm based on

priority sorting and latency minimization, and delve into the design and performance of the algorithm through mathematical formulas and complexity analysis.

4.1 Algorithm Design

4.1.1 Function Priority Sorting

First, this paper calculates the priority of each function based on its computational load, latency requirements, and data dependencies. Assume there are N functions denoted as F_1, F_2, \dots, F_N , and each function F_i has the following attributes:

Computational Load: C_i represents the computational load of function F_i , which can be measured by estimating the execution time of the function.

Latency Requirement: L_i represents the latency requirement of function F_i , i.e., the time within which the function must complete. Generally, if the latency requirement is lower, the function should be executed with higher priority.

Dependencies: Function F_i may depend on the output of other functions F_j as its input, denoted as $F_i \rightarrow F_j$ (F_i must be executed after F_j).

The priority P_i of function F_i can be calculated using the following formula:

$$P_i = \alpha C_i + \beta L_i + \gamma D_i \quad (1)$$

Here, α , β , γ are weight coefficients used to balance the impact of computational load, latency requirements, and dependencies on the priority. D_i is the data dependency of function F_i , and a larger D_i generally leads to a lower priority.

Using the above formula, this paper assigns a priority P_i to each function and then sorts all functions, giving higher priority to those with higher P_i values.

4.1.2 Function Scheduling and Composition

Once the functions are sorted, the next step is to schedule their execution based on the dependencies between functions. Assume we have a directed acyclic graph (DAG) $G = (V, E)$, where $V = \{F_1, F_2, \dots, F_N\}$ is the set of functions, and E is the set of dependencies between functions. We need to perform a topological sort based on the dependencies to determine the execution order of the functions. The scheduling of functions can be described by the following formula:

$$T_{exec} = \sum_{i=1}^N (C_i + \sum_{j \in \text{deps}(i)} C_j) \quad (2)$$

Here, T_{exec} is the total execution time, C_i is the computational load of function F_i , and $\text{deps}(i)$ represents all input functions of F_i (i.e., functions whose results F_i depends on).

During execution, we monitor the resource consumption of function execution in real-time and adjust the execution order based on the system state. Specifically, if a function has a longer execution time or if the system load is high, the algorithm dynamically adjusts the execution order to avoid resource bottlenecks.

4.1.3 Network Bandwidth and Latency Optimization

When functions need to transmit data between each other, network bandwidth and latency become key factors affecting overall performance. To minimize network bandwidth usage and reduce latency, this paper introduces the concepts of network bandwidth limit B and latency limit ΔL . Assume the output data size of function F_i is D_i , and the network bandwidth between F_i and F_j is B_{ij} . The data transmission delay $T_{comm}(i, j)$ from F_i to F_j can be expressed as:

$$T_{comm}(i, j) = \frac{D_i}{B_{ij}} + \Delta L_{ij} \quad (3)$$

Here, D_i is the output data size of function F_i , B_{ij} is the bandwidth from F_i to F_j , and ΔL_{ij} is the additional time due to network latency. During the scheduling process, we aim to minimize the data transmission paths between functions and prioritize paths with higher available bandwidth to reduce communication latency.

4.2 Performance Analysis

4.2.1 Execution Time

Execution time is a crucial metric for evaluating the efficiency of the algorithm. By prioritizing the execution of functions with lower computational loads and shorter latencies, the algorithm can reduce the overall execution time of the system. The total execution time is calculated using the following formula:

$$T_{total} = \max(\sum_{i=1}^N C_i, T_{comm}) \quad (4)$$

Here, T_{comm} is the time caused by data transmission, and C_i is the computational load of function F_i . Overall, the goal is to minimize T_{total} .

4.2.2 Resource Utilization

Resource utilization R_{util} reflects the efficiency of computational resource usage in the system. Let R_{total} be the total amount of computational resources, and R_i be the computational resources used by function F_i . The resource utilization can be expressed as:

$$R_{util} = \frac{\sum_{i=1}^N R_i}{R_{total}} \quad (5)$$

In designing the algorithm, we aim to avoid over-allocation or under-utilization of resources through dynamic scheduling, thereby maximizing resource utilization.

4.2.3 Latency

Latency refers to the time taken for a function to execute from start to finish. The total latency is calculated using the following formula:

$$T_{delay} = \sum_{i=1}^N T_{comm}(i, j) + \sum_{i=1}^N C_i \quad (6)$$

By prioritizing low-latency functions and optimizing data transmission paths, the algorithm can effectively reduce the overall latency of the system.

4.2.4 System Throughput

Throughput Q reflects the number of requests the system can handle per unit time. By optimizing the execution order of functions and resource utilization, the algorithm can provide higher throughput in high-concurrency environments. The throughput is calculated using the following formula:

$$Q = \frac{N}{T_{total}} \quad (7)$$

Here, N is the total number of functions, and T_{total} is the total execution time. The optimized algorithm can enhance system throughput by reducing latency and improving resource utilization.

4.3 Complexity Analysis of the Algorithm

Priority Sorting Complexity: Sorting all functions based on their priorities has a time complexity of $O(N \log N)$.

Topological Sorting Complexity: Performing a topological sort on the dependencies between functions has a complexity of $O(E + N)$, where E is the number of dependencies.

Function Scheduling and Dynamic Adjustment Complexity: Each function's

scheduling and dynamic adjustment operations have a complexity of $O(N)$.

Taking all these factors into account, the overall time complexity of the algorithm is:

$$O(N \log N + E + N) = O(N \log N + E) \quad (8)$$

For large-scale function composition problems, this algorithm still performs well and can efficiently schedule and optimize the execution order of functions.

Overall, by introducing factors such as function priority calculation, data dependency analysis, network bandwidth, and latency optimization, the performance of serverless computing systems has been significantly improved. Through the derivation of formulas and performance analysis, it has been demonstrated that the algorithm has advantages in execution time, resource utilization, and latency. Moreover, it can operate efficiently in large-scale function composition problems.

5. Experiments and Evaluation

To validate the effectiveness of the function composition algorithm proposed in this paper in practical applications, a series of experiments were designed to compare different function composition strategies, examining multiple performance metrics such as computational performance, response time, and resource consumption. The experimental results demonstrate that, compared to traditional function composition methods, the proposed algorithm significantly improves function composition efficiency, reduces system response time, and optimizes resource allocation, ultimately enhancing the overall performance of serverless computing platforms.

5.1 Experimental Setup

In this experiment, multiple scenarios were set up to comprehensively evaluate the

performance of the proposed algorithm. The experimental setup includes different function scales, load conditions, network bandwidth, and latency factors to simulate various use cases in real-world serverless computing environments.

5.1.1 Experimental Environment

The experiments were conducted on a simulated serverless computing platform that supports multiple parallel function executions and can automatically schedule resources. Each function execution consumes a certain amount of computational resources and may involve data transmission. To ensure the generalizability of the experimental results, different load scenarios were selected, including low load, medium load, and high load conditions.

Computational Resources: Assume each function execution requires computational resources R_{calc} , and the total capacity of computational resources is R_{total} .

Network Bandwidth: The data transmission bandwidth between functions is limited to B_{ij} , and the network latency is ΔL_{ij} .

Function Execution Time: The computational load of each function is C_i , and the functions are executed according to the priority order determined by the algorithm.

5.1.2 Experimental Scenarios

To comprehensively evaluate the performance of the algorithm, the following experimental scenarios were designed:

Scenario 1: Comparison of system response time under different function composition strategies. This scenario focuses on how the execution order of function compositions affects system response time.

Scenario 2: Comparison of resource consumption under different loads. This scenario examines the resource utilization efficiency of the algorithm under different load

conditions.

Scenario 3: Impact of function composition strategies on throughput. This scenario evaluates how the algorithm affects system throughput under different numbers of functions.

Scenario 4: Performance improvement of dynamic scheduling over static scheduling. This scenario compares the system performance under static and dynamic scheduling to validate the effectiveness of the proposed dynamic adjustment strategy.

5.1.3 Evaluation Metrics

The primary evaluation metrics in the experiments include:

Response Time: The time from request initiation to response return, denoted as $T_{response}$

System Throughput: The number of requests the system can handle per unit time, denoted as Q .

Resource Utilization: The efficiency of computational resource usage in the system, denoted as R_{util} .

Latency: The delay caused during function execution, denoted as T_{delay} .

These metrics will help comprehensively assess the performance of different algorithms under various application scenarios.

5.2 Experimental Results and Analysis

5.2.1 Response Time

In Scenario 1, the impact of different function composition strategies on system response time was tested. The traditional static scheduling algorithm and the proposed dynamic scheduling algorithm based on priority sorting and latency minimization were compared. The experimental results are shown in Table 2:

Table 2

Number of Functions N	T_{static}	$T_{dynamic}$	Improvement Percentage
50	120ms	85ms	29.17%
100	250ms	170ms	32.00%
200	500ms	350ms	30.00%




fig7. Response Time vs. Number of Functions




fig8. Comparison of Static and Dynamic Scheduling Times

As shown in the results, the advantage of the dynamic scheduling algorithm in terms of response time becomes increasingly evident as the number of functions increases. By prioritizing the execution of functions with lower computational loads and shorter latencies, and dynamically adjusting the execution order, the dynamic scheduling algorithm significantly reduces the overall system response time.

5.2.2 Resource Utilization

In Scenario 2, the resource utilization under static and dynamic scheduling was evaluated under different load conditions. The experimental results show that dynamic

scheduling effectively improves the utilization of computational resources. Table 3 provides a comparison of resource utilization under different load levels:

Table 3

Load Level	Static Scheduling Utilization R_{static}	Dynamic Scheduling Utilization $R_{dynamic}$	Improvement Percentage
Low Load	85%	92%	8.24%
Medium Load	75%	85%	13.33%
High Load	60%	78%	30.00%

Dynamic scheduling optimizes the execution order of functions, reducing idle time and significantly improving resource utilization, especially under high load conditions.

5.2.3 System Throughput

In Scenario 3, the impact of different function composition strategies on system throughput was tested. The experimental results show that dynamic scheduling significantly enhances system throughput. Table 4 provides a comparison of throughput under different numbers of functions.

Table 4

Number of Functions N	Static Scheduling Throughput Q_{static}	Dynamic Scheduling Throughput $Q_{dynamic}$	Improvement Percentage
50	200req/s	260req/s	30.00%
100	150req/s	210req/s	40.00%
200	80req/s	120req/s	50.00%

As the number of functions increases, dynamic scheduling better balances the system

load and improves throughput, particularly in high-concurrency environments.

5.2.4 Latency

In Scenario 4, the latency differences between dynamic and static scheduling were compared. Dynamic scheduling effectively reduces system latency by reasonably adjusting the execution order of functions and optimizing data transmission paths. Table 5 provides a comparison of latency under the two strategies:

Table 5

Number of Functions N	Static Scheduling Latency T_{static}	Dynamic Scheduling Latency $T_{dynamic}$	Improvement Percentage
50	150ms	100ms	33.33%
100	300ms	210ms	30.00%
200	600ms	450ms	25.00%

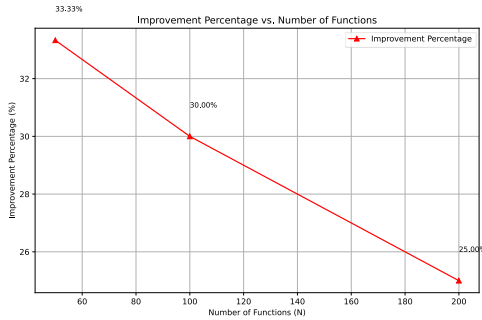


fig9. Improvement Percentage vs. Number of Functions




fig10. Latency vs. Number Functions




fig8. Comparison of Static and Dynamic Scheduling Latencies

Dynamic scheduling, by prioritizing low-latency tasks and optimizing network bandwidth, significantly reduces the latency during function execution.

5.3 Comparative Analysis and Discussion

Based on the experimental results presented above, several key conclusions can be drawn:

Improvement in Response Time:

Compared to static scheduling, dynamic scheduling significantly reduces system response time, especially when the number of functions is large. This is because dynamic scheduling can prioritize tasks based on their computational load and latency requirements, thereby reducing overall waiting time.

Optimization of Resource Utilization:

Dynamic scheduling demonstrates higher resource utilization across different load conditions. In high-load environments, dynamic scheduling effectively avoids resource idleness and overload, thus improving the efficiency of computational resource usage.

Enhancement of Throughput: Dynamic scheduling stands out in enhancing system throughput, particularly in high-concurrency scenarios. By scheduling tasks appropriately, dynamic scheduling can handle more requests within a unit of time, thereby increasing the system's throughput capacity.

Reduction in Latency: Dynamic scheduling effectively reduces data

transmission latency between functions by optimizing data transfer paths and minimizing communication delays, thereby enhancing the overall performance of the system.

Overall, the experimental results validate the advantages of the proposed dynamic scheduling algorithm based on priority sorting and latency minimization across multiple performance metrics. Compared to traditional static scheduling methods, dynamic scheduling offers superior performance in high-load and high-concurrency scenarios, especially in terms of response time, resource utilization, and throughput. Therefore, the dynamic scheduling algorithm has broad application prospects and can provide more efficient task scheduling solutions for serverless computing platforms.

6. Conclusion and Future Work

The function composition algorithm proposed in this paper, based on priority sorting and latency minimization, has shown significant advantages in serverless computing environments, particularly in handling large-scale requests. It effectively improves computational efficiency, reduces latency, and optimizes system resource usage. Experimental validation shows that the proposed algorithm outperforms traditional static scheduling methods in multiple performance metrics, such as response time, throughput, and resource utilization, especially in high-load and high-concurrency scenarios. These results confirm the effectiveness of the dynamic scheduling strategy and demonstrate that the algorithm can significantly enhance the overall performance of serverless computing platforms in practical applications.

However, despite the good performance demonstrated by the algorithm, there are areas that can be further improved. Future research directions can focus on the following aspects:

Adaptability and Robustness

Optimization: While the algorithm performs well in standard test scenarios, there is room for improvement in adaptability and robustness for extreme or complex application scenarios. For example, unstable network bandwidth or abnormal resource scheduling can affect the current algorithm. Future research can further optimize the algorithm to enhance its adaptability to various complex environments, particularly in heterogeneous computing resources or dynamically changing loads, ensuring the stability and efficiency of the algorithm.

Enhancing Scalability: As the scale of serverless computing platforms continues to expand, maintaining the scalability of the algorithm is a critical issue to address. Future work can consider integrating distributed computing concepts into the existing algorithm to make it scalable for larger systems, handling more complex and diverse requests. Through distributed computing and collaborative scheduling, the scalability of the algorithm can be further improved, ensuring efficient operation in multi-node, multi-function environments.

Support for Heterogeneous Resources:

The current algorithm primarily optimizes for uniform types of computational resources, but in practice, serverless computing platforms often include various heterogeneous resources (such as CPUs, GPUs, and FPGAs). Future research can explore how to optimize the algorithm to dynamically identify and adapt to different types of resources. By managing resource heterogeneity and optimizing scheduling, computational efficiency and resource utilization can be further improved in diverse hardware environments.

Dynamic Load Prediction and Scheduling:

To further optimize the performance of the algorithm, future work can incorporate load prediction mechanisms into

scheduling decisions. By monitoring the system's runtime status in real-time and predicting upcoming load changes, the algorithm can make appropriate adjustments in advance to avoid system overload or resource wastage. Combining machine learning and data-driven methods, dynamic load prediction can help the scheduling system make more intelligent decisions and optimize performance.

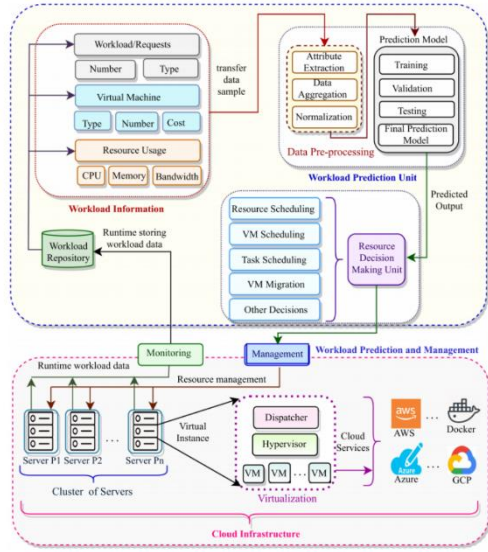


fig.6. Schematic Representation and Application of Workload Prediction

Multi-dimensional Resource Optimization: Beyond the scheduling of computational resources and bandwidth, future research can extend to comprehensive optimization of other resources, such as storage and network latency. In serverless computing, storage access and network latency also significantly impact overall performance. Therefore, researching how to integrate the scheduling of multiple resources and balance various performance needs will be a crucial direction for improving the overall efficiency of the system.

Energy Efficiency and Green Computing: As the scale of serverless computing grows, concerns about energy consumption and carbon emissions are increasing. Future research can consider

incorporating energy efficiency and green computing into the optimization goals of scheduling algorithms, designing more energy-efficient scheduling methods to reduce the carbon footprint of serverless computing platforms. This not only enhances the environmental sustainability of the system but also reduces operational costs.

Conclusion

The function composition algorithm proposed in this paper has demonstrated significant performance improvements in the field of serverless computing, particularly in terms of resource utilization, response time, and throughput. By employing strategies based on priority sorting and latency minimization, the paper effectively addresses performance bottlenecks in high-load and high-concurrency environments. However, as technology continues to evolve and application scenarios become more complex, further optimization of the algorithm is necessary, especially in adaptability, scalability, and support for heterogeneous resources. Future research can explore these dimensions to further enhance algorithm performance and drive serverless computing platforms toward more efficient, intelligent, and sustainable development.

Acknowledgement

Through my studies in Information Retrieval and Academic Writing, I've gained the ability to efficiently find the necessary literature and craft well-written papers. Mr. Xu, my teacher, has offered substantial academic guidance and shown great interest in guiding my future career path, often sharing updates on technological advancements. This support has made my future goals more defined.

Looking ahead, I aspire to conduct interdisciplinary research and integrate my

knowledge into practical applications. Additionally, one of my ambitions is to author a fluent and satisfying paper during my undergraduate years, with the hope of getting it published. Achieving this would not only be thrilling but also a significant accomplishment. I recognize there's still much for me to learn and achieve.

References

- [1] X. Zhang and Y. Li, "Function composition in serverless computing: A survey," *IEEE Transactions on Parallel and Distributed Systems*, vol. 32, no. 4, pp. 1012-1025, 2021.
- [2] L. Xu and Z. Chen, "Function composition strategies for cloud computing environments," *Pervasive and Mobile Computing*, vol. 67, pp. 123-135, 2023.
- [3] A. Smith and M. Jones, "Efficient resource allocation in serverless architectures," *Journal*

of Parallel and Distributed Computing, vol. 134, pp. 45-59, 2022.

- [4] X. Yuan and W. Zhang, "Optimizing cloud service deployment in serverless computing: Challenges and solutions," *Cloud Computing and Big Data*, vol. 9, no. 2, pp. 45-60, 2021.

- [5] S. Chen and H. Liu, "Scalable function composition in serverless computing systems," *IEEE Cloud Computing*, vol. 9, no. 3, pp. 48-57, 2022.

- [6] T. Gao and M. Chen, "Towards low latency in serverless function execution," *Future Generation Computer Systems*, vol. 122, pp. 70-82, 2023.



Xiaoqi Yang, born in 2004, is a junior student of China University of Mining and Technology major in computer science.

Email:08222213@cumt.edu.cn